

In [13]:

#This project predicts whether a customer will repay their credit within 90 days. The tasks and dataset were posted by Yuri Kashnitsky on Kaggle.

```
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
from matplotlib import rcParams
rcParams['figure.figsize'] = 7, 5
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier

#for each columns, replacing NaNs with the column median
def fill_nan(table):
    for col in table.columns:
        table[col] = table[col].fillna(table[col].median())
    return table

data = pd.read_csv('/Users/iqra/Desktop/credit_scoring_sample.csv')
data.head()
```

Out[13]:

	SeriousDlqin2yrs	age	NumberOfTime30-59DaysPastDueNotWorse	DebtRatio	NumberOfTimes90DaysLate	
0	0	64	0	0.249908	0	
1	0	58	0	3870.000000	0	
2	0	41	0	0.456127	0	
3	0	43	0	0.000190	0	
4	1	49	0	0.271820	0	

In [14]:

```


ax = data['SeriousDlqin2yrs'].hist(orientation='vertical', color='red')
ax.set_xlabel("number_of_observations")
ax.set_ylabel("unique_value")
ax.set_title("Target distribution")

print('Distribution of the target:')
data['SeriousDlqin2yrs'].value_counts()/data.shape[0]

independent_columns_names = [x for x in data if x != 'SeriousDlqin2yrs']
independent_columns_names
#visualizing the target distribution
table = fill_nan(data)
X = table[independent_columns_names]
y = table['SeriousDlqin2yrs']

```

Distribution of the target:



In [29]:

```

#Logistic Regression
lr = LogisticRegression(random_state=5, class_weight='balanced')
#finding the best regulation coefficient for our model
parameters = {'C': (0.0001, 0.001, 0.01, 0.1, 1, 10)}
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=5)
grid_search = GridSearchCV(lr, parameters, scoring='roc_auc', cv=skf)
grid_search = grid_search.fit(X, y)
grid_search.best_estimator_
grid_search.best_score_

lr = LogisticRegression(C=0.001, random_state=5, class_weight='balanced')
scal = StandardScaler()
lr.fit(scal.fit_transform(X), y)

```

Out[29]:

LogisticRegression(C=0.001, class_weight='balanced', random_state=5)

In [40]:

```
pd.DataFrame({'feat': independent_columns_names,
              'coef': lr.coef_.flatten().tolist()}), sort_values(by='coef', ascending=False)
#creating a random forest function
rf = RandomForestClassifier(n_estimators=100, n_jobs=-1, random_state=42,
                           class_weight='balanced')
#searching for the best params
parameters = {'max_features': [1, 2, 4], 'min_samples_leaf': [3, 5, 7, 9], 'max_depth': [5, 10, 15]}
rf_grid_search = GridSearchCV(rf, parameters, n_jobs=-1, scoring='roc_auc', cv=5, verbose=True)
rf_grid_search = rf_grid_search.fit(X, y)
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 53.4s

[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed: 5.6min finished

0.026820670654852385

In [25]:

```

#setting up params for bagging
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import cross_val_score, RandomizedSearchCV
parameters = {'max_features': [2, 3, 4], 'max_samples': [0.5, 0.7, 0.9],
              'base_estimator__C': [0.0001, 0.001, 0.01, 1, 10, 100]}

#we will use logistic regression for its ability to better interpret and evaluate features
bg = BaggingClassifier(LogisticRegression(class_weight='balanced'),
                      n_estimators=100, n_jobs=-1, random_state=42)
r_grid_search = RandomizedSearchCV(bg, parameters, n_jobs=-1,
                                   scoring='roc_auc', cv=skf, n_iter=20, random_state=1,
                                   verbose=True)
r_grid_search = r_grid_search.fit(X, y)

#finding the ROC AUC
r_grid_search.best_score_
r_grid_search.best_estimator_

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 5.2min
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 13.1min finished

```

Out[25]:

```

BaggingClassifier(base_estimator=LogisticRegression(C=100,
                                                    class_weight='balanced'),
                  max_features=2, max_samples=0.9, n_estimators=100,
                  n_jobs=-1,
                  random_state=42)

```

In []: